

# Development and Operation of Distributed Applications

Group 21

Ceylin Ece 5716950  
Dibyendu Gupta 5031125  
Taeyong Kwon 6283276  
Guotao Gou 6380743

# Automatic Versioning & Release

## Stable Release workflow:

1. Runs only on manual **release** or **workflow** to avoid conflict.
2. Remove the SNAPSHOT, then git tag this version, then publish the release.
3. Then increment the version and then append SNAPSHOT.

```
<groupId>tudelft.doda2025.team21</groupId>
<artifactId>lib-version</artifactId>
<version>1.5.1-SNAPSHOT</version>
```

Pom.xml -> source of truth

# Automatic Versioning & Release

## Pre Release workflow:

1. Runs only on every push.
2. Remove the SNAPSHOT, then set version as  
\${BASE\_VERSION}-\${BRANCH\_CLEAN}-\${TIMESTAMP}

<> Install 1/2: Add this to pom.xml:

```
<dependency>
  <groupId>tudelft.doda2025.team21</groupId>
  <artifactId>lib-version</artifactId>
  <version>1.4.3-test_branch-251121-181733</version>
</dependency>
```

# Automatic Versioning & Release

```
- name: Determine base version from pom.xml
  id: base_version
  run: |
    # remove -SNAPSHOT to get base version
    mvn versions:set -DremoveSnapshot
    BASE_VERSION=$(mvn -q -Dexpression=project.version -DforceStdout help:evaluate)
    echo "BASE_VERSION=$BASE_VERSION" >> $GITHUB_ENV
    echo "Base version: $BASE_VERSION"

- name: Create pre-release version string
  id: prerelease
  run: |
    # branch name like: feature/new-api -> feature-new-api
    BRANCH_RAW="${GITHUB_REF_NAME}"
    BRANCH_CLEAN=$(echo "$BRANCH_RAW" | tr '/' '-')
    TIMESTAMP=$(date +'%y%m%d-%H%M%S')

    PRE_VERSION="${BASE_VERSION}-${BRANCH_CLEAN}-${TIMESTAMP}"

    echo "Pre-release version: $PRE_VERSION"
    echo "PRE_VERSION=$PRE_VERSION" >> $GITHUB_ENV

- name: Set pom.xml to pre-release version
  run: |
    mvn -B versions:set -DnewVersion=${PRE_VERSION}
    mvn -B versions:commit
```

# Software Reuse in Library

- App depend on stable version lib-version
- App is compiled into **jar** in github workflow. So docker image can directly download it from github, which runs **faster** than local compiling.

```
<dependency>
  <groupId>tudelft.doda2025.team21</groupId>
  <artifactId>lib-version</artifactId>
  <version>1.5.0</version>
</dependency>

</dependencies>

<repositories>
  <repository>
    <id>github</id>
    <url>https://maven.pkg.github.com/doda2025-team21/lib-version</url>
  </repository>
</repositories>
```

# Software Reuse in Library

- App print out lib-version's **VersionUtil** in main

```
(base) ✘ ~/TUD_Q2/CS4295-DevOps/assignment/app/ [fix-lib-version] docker run -it --name app_test -p 8080:8080 ghcr.io/doda2025-team21/app:0.2.19
lib-version version: 1.5.0
  
:: Spring Boot ::          (v3.5.7)  
  
2026-01-21T16:16:37.674Z INFO 1 --- [           main] frontend.Main                  : Starting Main v0.2.19 using Java 25.0.1 with PID 1  
2026-01-21T16:16:37.676Z INFO 1 --- [           main] frontend.Main                  : No active profile set, falling back to 1 default profile  
2026-01-21T16:16:39.033Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port 8080 (http)  
2026-01-21T16:16:39.043Z INFO 1 --- [           main] o.apache.catalina.core.StandardService : Starting service [Tomcat]  
2026-01-21T16:16:39.044Z INFO 1 --- [           main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.48]  
2026-01-21T16:16:39.085Z INFO 1 --- [           main] o.a.c.c.C.[Tomcat].[localhost].[/] : Initializing Spring embedded WebApplicationContext  
2026-01-21T16:16:39.085Z INFO 1 --- [           main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 1 ms  
  
Working with MODEL_HOST="http://host.docker.internal:8081"  
2026-01-21T16:16:39.792Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port 8080 (http) with context path  
2026-01-21T16:16:39.903Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port 9090 (http)  
2026-01-21T16:16:39.904Z INFO 1 --- [           main] o.apache.catalina.core.StandardService : Starting service [Tomcat]  
2026-01-21T16:16:39.905Z INFO 1 --- [           main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.48]  
2026-01-21T16:16:39.910Z INFO 1 --- [           main] o.a.c.c.C.[Tomcat-1].[localhost].[/] : Initializing Spring embedded WebApplicationContext  
2026-01-21T16:16:39.910Z INFO 1 --- [           main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 1 ms  
2026-01-21T16:16:39.962Z INFO 1 --- [           main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 2 endpoints beneath base path '/actuator'  
2026-01-21T16:16:40.033Z INFO 1 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port 9090 (http) with context path
```

# Container & Orchestration

- Github workflow: use **buildx** to build multi-architecture.
- Use **.env** to configure the docker compose deployment.

```
services:  
  model-service:  
    image: "${MODEL_IMAGE}:${TAG_VERSION:-1.0.5-monitoring}"  
    container_name: backend  
    ports:  
      - "${MODEL_PORT:-8081}:${MODEL_PORT:-8081}"  
    environment:  
      MODEL_PORT: "${MODEL_PORT:-8081}"  
    volumes:  
      - ./model:/app/output  
    restart: unless-stopped  
  
  app:  
    image: "${APP_IMAGE}:${TAG_VERSION:-1.1.0}"  
    container_name: frontend  
    ports:  
      - "${APP_PORT:-8080}:${APP_PORT:-8080}"  
    environment:  
      APP_PORT: "${APP_PORT:-8080}"  
      MODEL_HOST: "http://model-service:${MODEL_PORT:-8081}"  
    depends_on:  
      - model-service  
    restart: unless-stopped
```

# **Setting up (Virtual) Infrastructure**

# Provisioning a Kubernetes Cluster

- Kubernetes base platform is reproducibly provisioned.
  - (Vagrant + VirtualBox + Ansible + Flannel CNI + MetalLB + Ingress + Dashboard + Istio)
- Identical cluster on every setup
  - When re-provisioned => same cluster with the same behaviour and capabilities



# What the Automation Sets Up

- Vagrant builds:
  - ctrl
  - NUM\_WORKERS node-\*
  - hosts on a *host-only network* (controller at 192.168.56.100, workers starting at 192.168.56.101/102/...).
- Our Vagrantfile generates a valid inventory.cfg for Ansible that contains all (and only) active nodes



Machines are:

- Numbered
- Disposable
- Redundant
- Reproducible

# **Setting up Software Environment**

# Idempotent System Configuration

- Ansible playbook uses several built-in modules to achieve an *idempotent* provisioning. Some examples include:
  - imports team SSH keys => key is added only if missing
  - disables swap and fstab entries => ensures SWAP stays off
  - loads overlay/br\_nfnetfilter => kernel modules ensured loaded
  - enables kubelet => service enabled + running



# **Setting up Kubernetes**

# Kubernetes Cluster Setup

## Kubernetes Cluster Initialization

- Initialize control plane => advertise address 192.168.56.100 and pod CIDR 10.244.0.0/16 (idempotent via /etc/kubernetes/admin.conf check)
- Distribute kubeconfig => copies kubeconfig to ~/.kube for vagrant and fetches ./kubeconfig for the host
- Join workers automatically => delegates kubeadm token create --print-join-command to the controller and joins each worker if kubelet.conf is missing

## Cluster Networking

- Flannel CNI is installed => enables pod-to-pod networking

# Kubernetes Platform Add-ons



## Platform add-ons

- **Helm** for reproducible add-on deployments
- **MetalLB** with pool **192.168.56.90-192.168.56.99** and readiness **waits**
- Cluster has a working HTTP **Ingress Controller** => ingress-nginx with class nginx and load balancer IP **192.168.56.90**
- Kubernetes **Dashboard** via Helm with an ingress => **dashboard.local**
- Cluster has a working **Istio Gateway** => Istio 1.25.2 with ingress gateway IP **192.168.56.91**

## Helps with

- Monitoring => Ingress, Dashboard
- Traffic Management => Istio

# Checking the cluster health from host

```
(base) ceylinece@Ceylins-MacBook-Pro operation % KUBECONFIG=~/kubeconfig kubectl get nodes -o wide
KUBECONFIG=~/kubeconfig kubectl get pods -A
NAME      STATUS  ROLES     AGE   VERSION  INTERNAL-IP    EXTERNAL-IP  OS-IMAGE           KERNEL-VERSION      CONTAINER-RUNTIME
ctrl      Ready   control-plane  15m  v1.32.4  192.168.56.100  <none>        Ubuntu 24.04.3 LTS  6.8.0-86-generic  containerd://1.7.24
node-1    Ready   <none>    14m  v1.32.4  192.168.56.101  <none>        Ubuntu 24.04.3 LTS  6.8.0-86-generic  containerd://1.7.24
node-2    Ready   <none>    14m  v1.32.4  192.168.56.102  <none>        Ubuntu 24.04.3 LTS  6.8.0-86-generic  containerd://1.7.24
NAMESPACE          NAME                               READY  STATUS    RESTARTS  AGE
ingress-nginx      ingress-nginx-controller-5b57587585-4m2hr  1/1   Running   0          13m
istio-system       istio-ingressgateway-7cff6d9468-t7xrw   1/1   Running   0          10m
istio-system       istiod-867f8c9f4f-7gr6c    1/1   Running   0          10m
kube-flannel       kube-flannel-ds-4bpt7    1/1   Running   0          14m
kube-flannel       kube-flannel-ds-lpgst    1/1   Running   0          14m
kube-flannel       kube-flannel-ds-wp49m    1/1   Running   0          15m
kube-system        coredns-668d6bf9bc-2ng5j   1/1   Running   0          15m
kube-system        coredns-668d6bf9bc-p59m7   1/1   Running   0          15m
kube-system        etcd-ctrl                1/1   Running   0          15m
kube-system        kube-apiserver-ctrl   1/1   Running   0          15m
kube-system        kube-controller-manager-ctrl 1/1   Running   0          15m
kube-system        kube-proxy-b67r2     1/1   Running   0          14m
kube-system        kube-proxy-rqjlb     1/1   Running   0          15m
kube-system        kube-proxy-wcf5s     1/1   Running   0          14m
kube-system        kube-scheduler-ctrl   1/1   Running   0          15m
kubernetes-dashboard  kubernetes-dashboard-api-5b57c4d774-mtsds  1/1   Running   0          8m6s
kubernetes-dashboard  kubernetes-dashboard-auth-78f9f9c85c-n7jkm  1/1   Running   0          8m6s
kubernetes-dashboard  kubernetes-dashboard-kong-9b49bb989-55282  1/1   Running   0          12m
kubernetes-dashboard  kubernetes-dashboard-metrics-scrapers-74d8cb664-fbk49 1/1   Running   0          12m
kubernetes-dashboard  kubernetes-dashboard-web-59b8766cc-7czm5   1/1   Running   0          12m
metallb-system      controller-bb5f47665-w2mp5    1/1   Running   0          14m
metallb-system      speaker-6d765      1/1   Running   0          14m
metallb-system      speaker-6xn24     1/1   Running   0          14m
metallb-system      speaker-dzr41     1/1   Running   0          14m
```

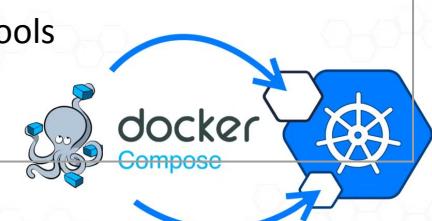
# Verifying load-balancers / ingress

```
(base) ceylinece@Ceylins-MacBook-Pro operation % KUBECONFIG=./kubeconfig kubectl get svc -n metallb-system
KUBECONFIG=./kubeconfig kubectl get svc -n ingress-nginx ingress-nginx-controller -o wide
KUBECONFIG=./kubeconfig kubectl get svc -n kubernetes-dashboard
KUBECONFIG=./kubeconfig kubectl get svc -n istio-system istio-ingressgateway -o wide
NAME          TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)         AGE
metallb-webhook-service   ClusterIP   10.109.229.130  <none>        443/TCP       15m
NAME          TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)         AGE   SELECTOR
ingress-nginx-controller   LoadBalancer  10.103.56.14   192.168.56.90  80:31662/TCP,443:30903/TCP  14m   app.kubernetes.io/component=controller,app.kubernetes.io/instance=ingress-nginx,app.kubernetes.io/name=ingress-nginx
NAME          TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)         AGE
kubernetes-dashboard-api   ClusterIP   10.105.177.118  <none>        8000/TCP     13m
kubernetes-dashboard-auth   ClusterIP   10.103.51.194  <none>        8000/TCP     13m
kubernetes-dashboard-kong-proxy   ClusterIP   10.110.242.179  <none>        443/TCP      13m
kubernetes-dashboard-metrics-scrapers   ClusterIP   10.99.56.243  <none>        8000/TCP     13m
kubernetes-dashboard-web   ClusterIP   10.96.141.36  <none>        8000/TCP     13m
NAME          TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)         AGE   SELECTOR
istio-ingressgateway   LoadBalancer  10.101.80.148   192.168.56.91  15021:31937/TCP,80:32693/TCP,443:31725/TCP  11m   app=istio-ingressgateway,istio=ingressgateway
```

# Kubernetes Usage

# Why are we moving from Docker-Compose →Kubernetes?

Docker Limitations	How Kubernetes addresses them?
Single-host constraint	Runs across cluster of machines; better for scalability and high availability
Basic Load-Balancing	Sophisticated load-balancing across pods in a cluster: <ul style="list-style-type: none"><li>- defining routing rules,</li><li>- configurable for canary and shadow releases</li><li>- DNS-based service discovery</li></ul>
No Self-healing	Kubernetes Scheduler and Controller Manager perform automatic container restarts, and rescheduling on pod failure.
App Monitoring	Allows exposing metrics that can be consumed by other tools (Prometheus) and perform visualizations (Grafana) and alerting (alertManager).





# How sms-checker Is Deployed on Kubernetes Clusters

- Deployed using Helm-chart
  - Package manager that describe Kubernetes resources
- Resources:
  - Deployments: Creates pods with specific labels.
  - Services: Allows for internal communication through Cluster IP.  
(accessible only within the cluster)
  - ServiceMonitor: Automatically discovers and scrapes metrics from K8 services.  
Bridges Prometheus and application.



Dashboard

Add

Dashboards

Display

Help

# php-apache

Namespace

## default

Condition

Summary All conditions OK

OK

Desired replicas

6

Up-to-date replicas

6

Ready replicas

5



Service: CRIT



Deployment problems

Pending: since 3 hours 12 minutes (warn/crit at 5 minutes 0 seconds/10 minutes 0 seconds)

Ready: 5/6, Up-to-date: 6/6, Not ready for: 3 hours 17 minutes (warn/crit at 5 minutes 0 seconds/10 minutes 0 seconds) CRIT

State

Service

Object

Name

CRIT

Status

Object

Pod

Name: php-apache-6cb4d4d74b-m6x8g

CRIT

Replicas

Object

Deployment

Name: php-apache

Deployment

CPU usage



CPU resources: Deployment

CPU usage

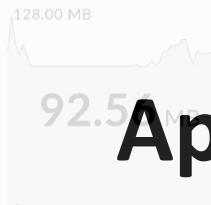


CPU usage: Pods

CPU usage



Memory usage



Memory resources: Deployment



Memory usage: Pods



# App Monitoring





Grafana

+



Prometheus

+



Alertmanager

# Monitoring the Cluster

## What are we monitoring?

- Kubernetes health
- App frontend metrics

## How?

- Endpoints are exposed through app-frontend endpoint /metrics
- Prometheus scrapes the Service for metrics via ServiceMonitors
- Grafana queries Prometheus to visualize data
- AlertManager receives alerts from Prometheus: when Prometheus rule is Firing
- AlertManager sends alerts to configured receivers (through E-mails, Slack, etc.)

## Relevance?

- Replicates production systems
- Infrastructure-as-code: managing and monitoring resources

# Alerting via Email

- Enabled and configured Alertmanager with SMTP-based email notifications
  - Triggers when traffic exceeds 2 requests/min for 1 minutes
- Injected SMTP credentials securely at install time via Helm values ; possible improvement using “Secrets”
- Verified alert lifecycle:  
Inactive → Pending → Firing → Resolved
- Integrated Grafana dashboards for real-time metric visualization

```
kind: PrometheusRule
metadata:
  name: high-traffic-alert
  labels:
    release: {{ .Release.Name }}
spec:
  groups:
    - name: traffic-alerts
      rules:
        - alert: HighRequestRate
          expr: sum(rate(sms_requests_total[1m])) * 60 > 2
          for: 1m
          labels:
            severity: warning
```

me

[ALERT] HighRequestRate - app - 1 alert for alertname=HighRequestRate View in Alertmanager [1] Resolved Labels alertname = HighRequestRate prometheus = ... 9:37 PM

me

[ALERT] HighRequestRate - app - 1 alert for alertname=HighRequestRate View in Alertmanager [1] Firing Labels alertname = HighRequestRate prometheus = defa... 9:32 PM



# Grafana Dashboard and Metrics

Metrics tracked: sms\_request\_total

<insert dashboard image>

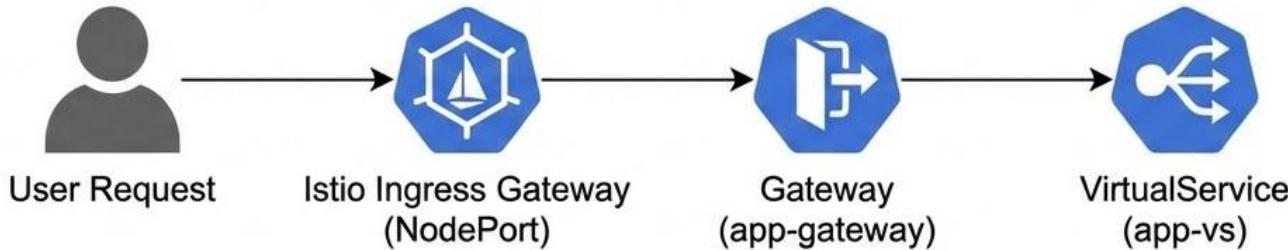
## Key deliverables of A4

Key deliverables:

**Traffic Management:** Gateway, VirtualService, and DestinationRules

- **Canary Release:** Weighted routing (90% Stable / 10% Canary)
- **Sticky Sessions:** Consistent hashing via cookies

# Traffic Management Architecture



## Routing Logic

- The VirtualService splits traffic based on headers (x-canary) or weights.
- Traffic is directed to specific DestinationRules which define the subsets: Stable: Version 0.2.17 (90% traffic). Canary: Version 0.2.14 (10% traffic).
- This ensures consistent routing between the app and model services (e.g., old-to-old, new-to-new).

# Canary Deployment Strategy

## Configuration:

- Defined in `values.yaml` under `istio.canary`
- **Default Route:** 90% weight to Stable, 10% weight to Canary
- **Header Override:** Requests with header `x-canary: true` are forced to the canary subset for testing purposes

**Note:** This fulfills the A4 requirement to "use DestinationRules for a canary release to a small fraction of the users" + 90/10 traffic

```
taea@taea-HP-Laptop-15-fd0xxx:~/doda-assignment/operation$  
for i in {1..10}; do  
  curl -s http://sms-istio.local/sms/ -X POST -H "Content-Type: application/json" -d '{"sms": "test '$i'"}' > /dev/null  
  sleep 1  
done  
  
echo "==== STABLE ===" && KUBECONFIG=./kubeconfig kubectl logs deploy/app-stable -c app --tail=20 | grep "test "  
echo "==== CANARY ===" && KUBECONFIG=./kubeconfig kubectl logs deploy/app-canary -c app --tail=20 | grep "test "  
==== STABLE ===  
Found 2 pods, using pod/app-stable-8447559955-p54mp  
Requesting prediction for "test message 1" ...  
Requesting prediction for "test message 2" ...  
Requesting prediction for "test message 5" ...  
Requesting prediction for "test 4" ...  
Requesting prediction for "test 5" ...  
Requesting prediction for "test 4" ...  
==== CANARY ===  
Requesting prediction for "test 6" ...
```

# Sticky Sessions (Stability)

**The Problem:** Ensuring users stay on the same version during an experiment session.

**The Solution:** HTTP Cookie-based consistent hashing.

**Implementation:**

- Cookie Name: sms-session.
- TTL: 3600 seconds (1 hour).

**How it works:**

1. First request: Istio routes based on the 90/10 weight.
2. Response: Includes Set-Cookie: sms-session=....
3. Subsequent requests: The cookie forces the user to the same subset (Stable or Canary).

```
● taea@taea-HP-Laptop-15-fd0xxx:~/doda-assignment/operation$ curl -c cookies.txt -b cookies.txt http://sms-istio.local/sms/ -d "message=test" {"timestamp": "2026-01-25T01:35:34.961+00:00", "status": 415, "error": "Unsupported Media Type", "path": "/sms/"}
● a-assignment/operation$ curl -c cookies.txt -b cookies.txt http://sms-istio.local/sms/ -d "message=test" {"timestamp": "2026-01-25T01:35:48.639+00:00", "status": 415, "error": "Unsupported Media Type", "path": "/sms/"}
● a-assignment/operation$ # Test sticky sessions - same cookie should hit same pod
curl -c cookies.txt -b cookies.txt http://sms-istio.local/sms/ -X POST -H "Content-Type: application/json" -d '{"sms": "sticky test 1"}'
curl -c cookies.txt -b cookies.txt http://sms-istio.local/sms/ -X POST -H "Content-Type: application/json" -d '{"sms": "sticky test 2"}'
curl -c cookies.txt -b cookies.txt http://sms-istio.local/sms/ -X POST -H "Content-Type: application/json" -d '{"sms": "sticky test 3"}'

# Check - all 3 should go to the SAME version
echo "==== STABLE ====" && KUBECONFIG=./kubeconfig kubectl logs deploy/app-stable -c app --tail=10 | grep "sticky"
echo "==== CANARY ====" && KUBECONFIG=./kubeconfig kubectl logs deploy/app-canary -c app --tail=10 | grep "sticky"
{"classifier": null, "result": "ham", "sms": "sticky test 1", "guess": null}, {"classifier": null, "result": "ham", "sms": "sticky test 2", "guess": null}, {"classifier": null, "result": "ham", "sms": "sticky test 3", "guess": null}
==== STABLE ====
Found 2 pods, using pod/app-stable-8447559955-p54mp
==== CANARY ====
Requesting prediction for "sticky" test 3" ...
```

# Verification: Rate Limiting Enforcement

## Normal Traffic ✓

Requests 1-10 (Within Burst Limit)

Request 1: HTTP 200 OK

Request 2: HTTP 200 OK

Request 3: HTTP 200 OK

...

Request 10: HTTP 200 OK



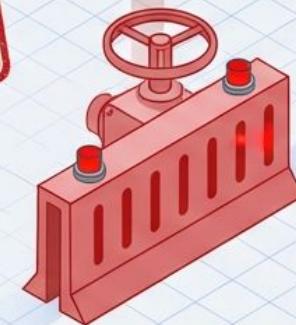
## Blocked Traffic ✗

Request 11 (Limit Exceeded)



## HTTP 429 Too Many Requests

x-local-rate-limit: TOO\_MANY\_REQUESTS



# Continuous Experimentation

- **The Experiment:** Testing a new frontend UI design in the canary version (0.2.14)
- **Hypothesis:** The new UI will improve user engagement
- **Observability:**
  - **Grafana Dashboard:** `decision-making-a4.json` performs an A/B comparison
  - **Metrics Tracked:** Latency distribution, Request Rate, and Error rate
- **Decision Process:** Compare stable vs. canary metrics in the dashboard to accept or reject the release.

# Configuration & Deployment

## Helm Integration:

- All Istio features are configurable via `values.yaml` (e.g.`istio.enabled`,  
`canary.enabled`, `stickySession.enabled`)
  - **Ingress Name:** Configurable via `ingressGateway.selector` to support different clusters, meeting the assignment requirement for portability
- 
- **Deployment Commands:**
    - Namespace labeled for Istio injection: `kubectl label namespace default istio-injection=enabled`
    - Deployed via Helm upgrade command

## Want to understand the flow completely?

[docs/deployment.md](#): Explains deployment structure and data flow

[docs/extension.md](#): Extension proposal for release engineering improvements

[docs/continuous-experimentation.md](#): Details the hypothesis and decision metrics